



# OPERATING SYSTEMS

## Synchronization Examples

---

**Suresh Jamadagni**

Department of Computer Science

# OPERATING SYSTEMS

## Slides Credits for all the PPTs of this course

---



- The slides/diagrams in this course are an **adaptation**, **combination**, and **enhancement** of material from the following resources and persons:
  1. Slides of Operating System Concepts, Abraham Silberschatz, Peter Baer Galvin, Greg Gagne - 9<sup>th</sup> edition 2013 and some slides from 10<sup>th</sup> edition 2018
  2. Some conceptual text and diagram from Operating Systems - Internals and Design Principles, William Stallings, 9<sup>th</sup> edition 2018
  3. Some presentation transcripts from A. Frank – P. Weisberg
  4. Some conceptual text from Operating Systems: Three Easy Pieces, Remzi Arpaci-Dusseau, Andrea Arpaci Dusseau

# OPERATING SYSTEMS

---

## Synchronization Examples


- ? Windows
- ? Linux
- ? Pthreads
- ? Solaris

**Suresh Jamadagni**

Department of Computer Science

- ❑ Uses interrupt masks to protect access to global resources on uniprocessor systems
- ❑ Uses **spinlocks** on multiprocessor systems
  - ❑ For reasons of efficiency, kernel ensures that a thread will never be preempted while holding a spinlock.
- ❑ Also provides **dispatcher objects** outside the kernel, to synchronize mutex locks, semaphores, events, and timers
  - ❑ **Events**
    - ▶ An event acts much like a condition variable (i.e notify a waiting thread when a desired condition occurs)
  - ❑ Timers notify one or more thread when time expired

- ❑ Dispatcher objects may be in either a **signaled-state** (object available and a thread will not block) or a **non-signaled state** (object not available, thread will block)
- ❑ A Relationship exists between the state of a dispatcher object and the state of a thread.
  - ❑ State of a thread changes from ready to waiting and vice-versa



[?] Linux:

[?] Prior to kernel Version 2.6, disables interrupts to implement short critical sections

[?] Version 2.6 and later, fully preemptive

[?] Linux provides:

[?] Semaphores

[?] atomic integers

[?] spinlocks

[?] reader-writer versions of both

[?] On single-cpu system, spinlocks replaced by enabling and disabling kernel preemption

- Atomic variables  
**atomic\_t** is the data type for atomic integer
- Consider the variables  
**atomic\_t counter;**  
**int value;**

<i>Atomic Operation</i>	<i>Effect</i>
<code>atomic_set(&amp;counter,5);</code>	<code>counter = 5</code>
<code>atomic_add(10,&amp;counter);</code>	<code>counter = counter + 10</code>
<code>atomic_sub(4,&amp;counter);</code>	<code>counter = counter - 4</code>
<code>atomic_inc(&amp;counter);</code>	<code>counter = counter + 1</code>
<code>value = atomic_read(&amp;counter);</code>	<code>value = 12</code>

- ❑ Pthreads API is OS-independent, widely used on UNIX, Linux, and macOS
- ❑ It provides:
  - ❑ mutex locks
  - ❑ semaphores
  - ❑ condition variable
- ❑ Non-portable extensions include:
  - ❑ read-write locks
  - ❑ spinlocks



- Creating and initializing the lock

```
#include <pthread.h>

pthread_mutex_t mutex;

/* create and initialize the mutex lock */
pthread_mutex_init(&mutex, NULL);
```

- Acquiring

```
/* acquire the mutex lock */
pthread_mutex_lock(&mutex);

/* critical section */

/* release the mutex lock */
pthread_mutex_unlock(&mutex);
```

- POSIX provides two versions – **named** and **unnamed**.
- Named semaphores (have actual names in the file system) can be shared by multiple unrelated processes
- Unnamed semaphores can be used only by threads belonging to the same process.

- Creating and initializing the semaphore:

```
#include <semaphore.h>
sem_t *sem;

/* Create the semaphore and initialize it to 1 */
sem = sem_open("SEM", O_CREAT, 0666, 1);
```

- Another process can access the semaphore by referring to its name **SEM**.

- Acquiring and releasing the semaphore:

```
/* acquire the semaphore */
sem_wait(sem);

/* critical section */

/* release the semaphore */
sem_post(sem);
```

- Creating and initializing the semaphore:

```
#include <semaphore.h>
sem_t sem;

/* Create the semaphore and initialize it to 1 */
sem_init(&sem, 0, 1);
```

- Acquiring and releasing the semaphore:

```
/* acquire the semaphore */
sem_wait(&sem);

/* critical section */

/* release the semaphore */
sem_post(&sem);
```

- Since POSIX is typically used in C/C++ and these languages do not provide a monitor (A high-level abstraction that provides a convenient and effective mechanism for process synchronization) , POSIX condition variables are associated with a POSIX mutex lock to provide mutual exclusion: Creating and initializing the condition variable:

```
pthread_mutex_t mutex;  
pthread_cond_t cond_var;  
  
pthread_mutex_init(&mutex, NULL);  
pthread_cond_init(&cond_var, NULL);
```

- Thread waiting for the condition **a == b** to become true:

```
pthread_mutex_lock(&mutex);  
while (a != b)  
    pthread_cond_wait(&cond_var, &mutex);  
  
pthread_mutex_unlock(&mutex);
```

- Thread signaling another thread waiting on the condition variable:

```
pthread_mutex_lock(&mutex);  
a = b;  
pthread_cond_signal(&cond_var);  
pthread_mutex_unlock(&mutex);
```

- ❑ Implements a variety of locks to support multitasking, multithreading (including real-time threads), and multiprocessing
- ❑ Uses **adaptive mutexes** for efficiency when protecting data from short code segments (< a few 100 instructions)
  - ❑ Starts as a standard semaphore spin-lock
  - ❑ If lock held, and by a thread running on another CPU, spins
  - ❑ If lock held by non-run-state thread (i.e. the thread holding the lock is not currently in run state), block and sleep waiting for signal of lock being released

- ❑ Uses **condition variables**
- ❑ Uses **readers-writers** locks when longer sections of code need access to data
- ❑ Uses **turnstiles** to order the list of threads waiting to acquire either an adaptive mutex or reader-writer lock
  - ❑ Turnstiles are per-lock-holding-thread, not per-object
- ❑ Priority-inheritance per-turnstile gives the running thread the highest of the priorities of the threads in its turnstile





**THANK YOU**

---

**Suresh Jamadagni**

Department of Computer Science and Engineering

**[sureshjamadagni@pes.edu](mailto:sureshjamadagni@pes.edu)**