

OPERATING SYSTEMS

Storage Management

Chandravva Hebbi

Department of Computer Science

OPERATING SYSTEMS

Mass-Storage Structure – Disk Scheduling
FCFS, SSTF, SCAN, C-SCAN, LOOK

Chandravva Hebbi

Department of Computer Science

- The slides/diagrams in this course are an **adaptation**, **combination**, and **enhancement** of material from the following resources and persons:
 1. Slides of Operating System Concepts, Abraham Silberschatz, Peter Baer Galvin, Greg Gagne - 9th edition 2013 and some slides from 10th edition 2018
 2. Some conceptual text and diagram from Operating Systems - Internals and Design Principles, William Stallings, 9th edition 2018
 3. Some presentation transcripts from A. Frank – P. Weisberg
 4. Some conceptual text from Operating Systems: Three Easy Pieces, Remzi Arpaci-Dusseau, Andrea Arpaci Dusseau

- ❑ The operating system is responsible for using hardware efficiently — for the disk drives, this means having a fast access time and disk bandwidth
- ❑ Minimize seek time
- ❑ Seek time \approx seek distance
- ❑ Disk **bandwidth** is the total number of bytes transferred, divided by the total time between the first request for service and the completion of the last transfer

- ❑ There are many sources of disk I/O request
 - ❑ OS
 - ❑ System processes
 - ❑ Users processes
- ❑ I/O request includes input or output mode, disk address, memory address, number of sectors to transfer
- ❑ OS maintains queue of requests, per disk or device
- ❑ Idle disk can immediately work on I/O request, busy disk means work must queue
 - ❑ Optimization algorithms only make sense when a queue exists

- ❑ Note that drive controllers have small buffers and can manage a queue of I/O requests (of varying “depth”)
- ❑ Several algorithms exist to schedule the servicing of disk I/O requests
- ❑ The analysis is true for one or many platters
- ❑ We illustrate scheduling algorithms with a request queue (0-199)


98, 183, 37, 122, 14, 124, 65, 67

Head pointer 53

Illustration shows total head movement of 640 cylinders

queue = 98, 183, 37, 122, 14, 124, 65, 67


head starts at 53



OPERATING SYSTEMS

SSTF

- ? Shortest Seek Time First selects the request with the minimum seek time from the current head position
- ? SSTF scheduling is a form of SJF scheduling; may cause starvation of some requests
- ? Illustration shows total head movement of 236 cylinders




$$12+2+30+23+84+24+2+59 = 236 \text{ cylinders}$$

- ❑ The disk arm starts at one end of the disk, and moves toward the other end, servicing requests until it gets to the other end of the disk, where the head movement is reversed and servicing continues.
- ❑ **SCAN algorithm** Sometimes called the **elevator algorithm**
- ❑ But note that if requests are uniformly dense, largest density at other end of disk and those wait the longest

queue = 98, 183, 37, 122, 14, 124, 65, 67

head starts at 53



Number of cylinder moves=
 $16+23+79+2+31+24+2+59=236$


- ❑ Provides a more uniform wait time than SCAN
- ❑ The head moves from one end of the disk to the other, servicing requests as it goes
 - ❑ When it reaches the other end, however, it immediately returns to the beginning of the disk, without servicing any requests on the return trip
- ❑ Treats the cylinders as a circular list that wraps around from the last cylinder to the first one
- ❑ Total number of cylinders?

OPERATING SYSTEMS

C-SCAN (Cont.)

queue = 98, 183, 37, 122, 14, 124, 65, 67

head starts at 53




$$\begin{aligned} &= (65-53) + (67-65) + (98-67) + (122-98) + (124-122) + (183-124) + (199-183) + (199-0) + (14-0) + (37-14) \\ &= 12 + 2 + 31 + 24 + 2 + 59 + 16 + 199 + 14 + 23 \\ &= 382 \end{aligned}$$

- ❑ LOOK a version of SCAN, C-LOOK a version of C-SCAN
- ❑ Arm only goes as far as the last request in each direction, then reverses direction immediately, without first going all the way to the end of the disk
- ❑ Total number of cylinders?

OPERATING SYSTEMS

C-LOOK (Cont.)



Total head movements incurred while servicing these requests

$$\begin{aligned} &= (65 - 53) + (67 - 65) + (98 - 67) + (122 - 98) + (124 - 122) + (183 - 124) + (183 - 14) + (37 - 14) \\ &= 12 + 2 + 31 + 24 + 2 + 59 + 169 + 23 \\ &= 322 \end{aligned}$$

- ❑ SSTF is common and has a natural appeal
- ❑ SCAN and C-SCAN perform better for systems that place a heavy load on the disk
 - ❑ Less starvation
- ❑ Performance depends on the number and types of requests
- ❑ Requests for disk service can be influenced by the file-allocation method
 - ❑ And metadata layout
- ❑ The disk-scheduling algorithm should be written as a separate module of the operating system, allowing it to be replaced with a different algorithm if necessary
- ❑ Either SSTF or LOOK is a reasonable choice for the default algorithm

? What about rotational latency?

- ? Difficult for OS to calculate
- ? The rotational latency can be nearly as large as the average seek time.
- ? It is difficult for the operating system to schedule for improved rotational latency, though, because modern disks do not disclose the physical location of logical blocks.
- ? Disk manufacturers have been alleviating this problem by implementing disk-scheduling algorithms in the controller hardware built into the disk drive.

? How does disk-based queueing effect OS queue ordering efforts?

- ? If the OS sends a batch of requests to the controller, the controller can queue them and then schedule them to improve both the seek time and the rotational latency.



THANK YOU

Chandravva Hebbi

Department of Computer Science Engineering

chandravvahebbi@pes.edu